

1. Basic blog articles
2. Serious tutorials
3. Basic application in NLP
 - 3.1. RL for structured prediction
 - 3.2. Learning to search (exposure bias)
 - 3.3. Misc.

1. Basic blog articles

- [Deep Reinforcement Learning: Pong from Pixels](#), Andrej Karpathy's blog post on DRL.
- [Deep Reinforcement Learning Doesn't Work Yet](#), undergraduate now at Google.

2. Serious tutorials

- [Deep Reinforcement Learning, Decision Making, and Control](#), ICML 2017 tutorials, Berkeley.
- [Deep Reinforcement Learning through Policy Optimization](#), with an emphasis on policy gradient/optimization, Berkeley.
- [Deep Reinforcement Learning](#), ICML 2016 tutorials, David Silver, DeepMind.

3. Basic application in NLP

3.1. RL for structured prediction

- [Reward augmented maximum likelihood for neural structured prediction](#), NIPS 2016.
- [SeqGAN: Sequence Generative Adversarial Nets with Policy Gradient](#), AAAI 2017.
- [An Actor-Critic Algorithm for Sequence Prediction](#), ICLR 2016.
- [Self-critical sequence training for image captioning](#), CVPR 2017.
- [Decoding with Value Networks for Neural Machine Translation](#), NIPS 2017.

3.2. Learning to search (exposure bias)

- [Sequence-to-sequence learning as beam-search optimization](#), EMNLP 2016.
- [Professor forcing: a new algorithm for training recurrent networks](#), NIPS 2016.
- [SEARNN: Training RNNs with global-local losses](#), ICLR 2018.

3.3. Misc.

- [Sequence Tutor: Conservative Fine-Tuning of Sequence Generation Models with KL-control](#), ICML 2017.
- [Seq2SQL: Generating Structured Queries From Natural Language Using Reinforcement Learning](#), ICLR 2018, **reject**.
- [Maximum Margin Reward Networks for Learning from Explicit and Implicit Supervision](#), EMNLP 2018.